

Genre Classification

Alexandria Simms

2025-07-14

Contents

1	Methods: Alexandria Simms	12
2	K Nearest Neighbors Model	12
3	Random Forest Model	13
4	Model Tuning	16
5	Results	22

```
features <- fread("fma_metadata/features.csv", data.table = FALSE) %>%
  janitor::clean_names()

tracks_raw <- read_csv("fma_metadata/tracks.csv", col_names = FALSE)
```

```
## Warning: One or more parsing issues, call 'problems()' on your data frame for details
## e.g.:
##   dat <- vroom(...)
##   problems(dat)
```

```
## Rows: 106577 Columns: 53
## -- Column specification -----
## Delimiter: ","
## chr (52): X2, X3, X4, X5, X6, X7, X8, X9, X10, X11, X12, X13, X14, X15, X16,...
## dbl (1): X1
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
header1 <- tracks_raw[1, ] %>% unlist(use.names = FALSE)
header2 <- tracks_raw[2, ] %>% unlist(use.names = FALSE)
column_names <- make.names(paste0(header1, ".", header2), unique = TRUE)
tracks <- tracks_raw[-c(1, 2), ]
colnames(tracks) <- column_names
tracks$track_id <- as.integer(rownames(tracks))
```

```
tracks_clean <- tracks %>%
  select(track_id, genre_top = `track.genre_top`, subset = `set.subset`) %>%
  filter(!is.na(genre_top), subset == "small") %>%
  mutate(genre_top = as.factor(genre_top))
```

```
colnames(features)[1] <- "track_id"
features <- features[-1, ]
features$track_id <- as.integer(features$track_id)
```

Warning: NAs introduced by coercion

```
combined_data <- inner_join(tracks_clean, features, by = "track_id")
```

```
colnames(features)[1] <- "track_id"
features <- features[-1, ]
features$track_id <- as.integer(features$track_id)
combined_data <- inner_join(tracks_clean, features, by = "track_id")
```

“{r. Normalize Features}

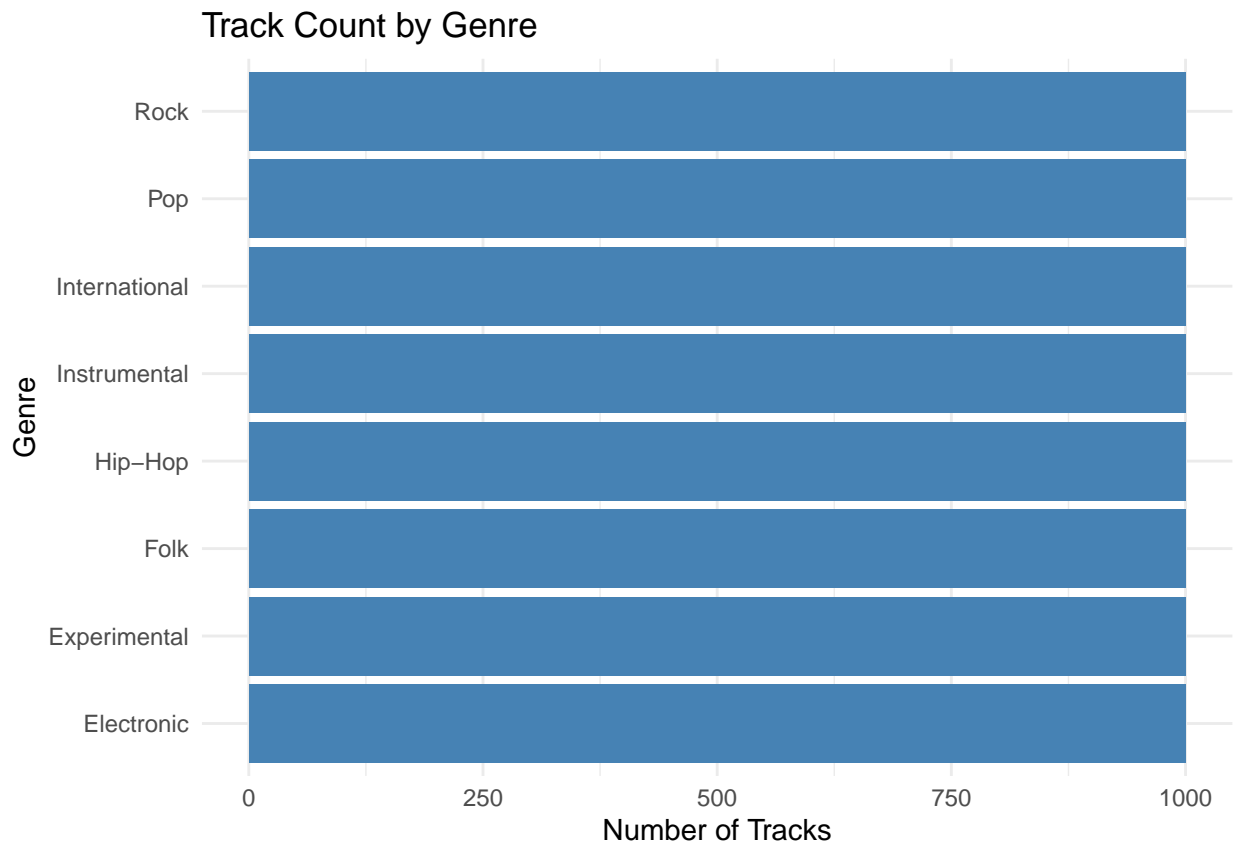
```
features <- features[-1, ] featuretrack_id <- as.integer(featuretrack_id) combined_data
<- inner_join(tracks_clean, features, by = "track_id") numeric_features <- com-
bined_data %>% select(-track_id, -subset, -genre_top) %>% select(where(is.numeric))
feature_matrix <- scale(numeric_features)
model_data <- data.frame(genre_top = combined_data$genre_top, feature_matrix)
```

\newpage

Exploratory Analysis Alexandria Simms

```
““ r
genre_counts <- tracks_clean %>%
  count(genre_top, sort = TRUE)
```

```
ggplot(genre_counts, aes(x = reorder(genre_top, n), y = n)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  coord_flip() +
  labs(title = "Track Count by Genre",
       x = "Genre", y = "Number of Tracks") +
  theme_minimal()
```



```
# Double-check subset distribution
table(tracks_clean$subset)
```

```
##
## small
## 8000
```

```
# Identify MFCC, chroma, and spectral columns
mfcc_cols <- grep("^mfcc", names(combined_data), value = TRUE)
chroma_cols <- grep("^chroma", names(combined_data), value = TRUE)
spectral_cols <- grep("^spectral", names(combined_data), value = TRUE)

# Convert relevant columns to numeric safely
```

```

combined_data <- combined_data %>%
  mutate(across(all_of(c(mfcc_cols, chroma_cols, spectral_cols)), ~as.numeric(.)))

# Calculate per-track means
audio_means <- combined_data %>%
  mutate(
    mfcc_mean = rowMeans(select(., all_of(mfcc_cols)), na.rm = TRUE),
    chroma_mean = rowMeans(select(., all_of(chroma_cols)), na.rm = TRUE),
    spectral_mean = rowMeans(select(., all_of(spectral_cols)), na.rm = TRUE)
  ) %>%
  select(genre_top, mfcc_mean, chroma_mean, spectral_mean)

# Make sure mfcc_cols is defined
mfcc_cols <- grep("^mfcc", colnames(combined_data), value = TRUE)

# Convert MFCC columns to numeric if needed
combined_data <- combined_data %>%
  mutate(across(all_of(mfcc_cols), ~as.numeric(.)))

# Create mean column
combined_data <- combined_data %>%
  mutate(mfcc_mean = rowMeans(select(., all_of(mfcc_cols)), na.rm = TRUE))

# Get column names that contain chroma
chroma_cols <- grep("^chroma", colnames(combined_data), value = TRUE)

# Convert to numeric and calculate the mean
combined_data <- combined_data %>%
  mutate(across(all_of(chroma_cols), ~as.numeric(.))) %>%
  mutate(chroma_mean = rowMeans(select(., all_of(chroma_cols)), na.rm = TRUE))

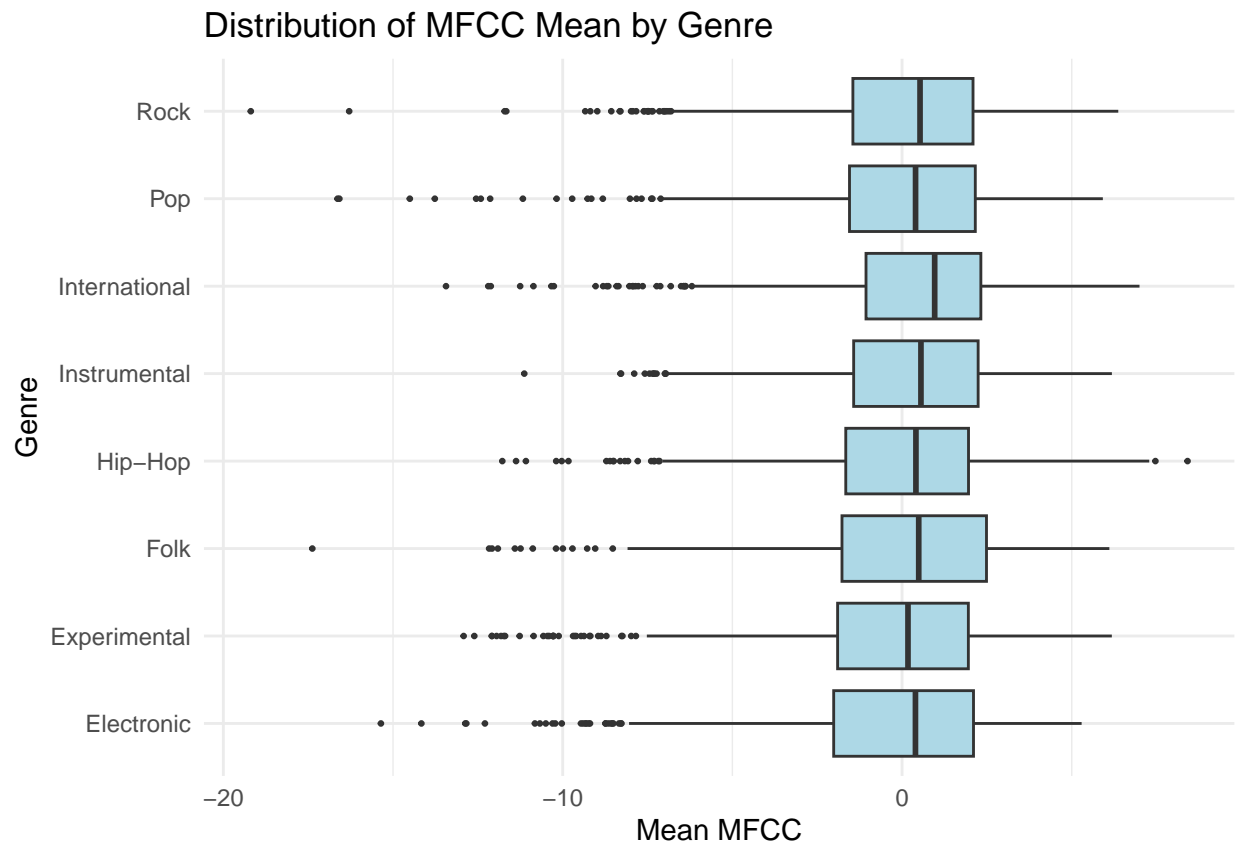
# Get column names that contain spectral
spectral_cols <- grep("^spectral", colnames(combined_data), value = TRUE)

# Convert to numeric and calculate the mean
combined_data <- combined_data %>%
  mutate(across(all_of(spectral_cols), ~as.numeric(.))) %>%
  mutate(spectral_mean = rowMeans(select(., all_of(spectral_cols)), na.rm = TRUE))

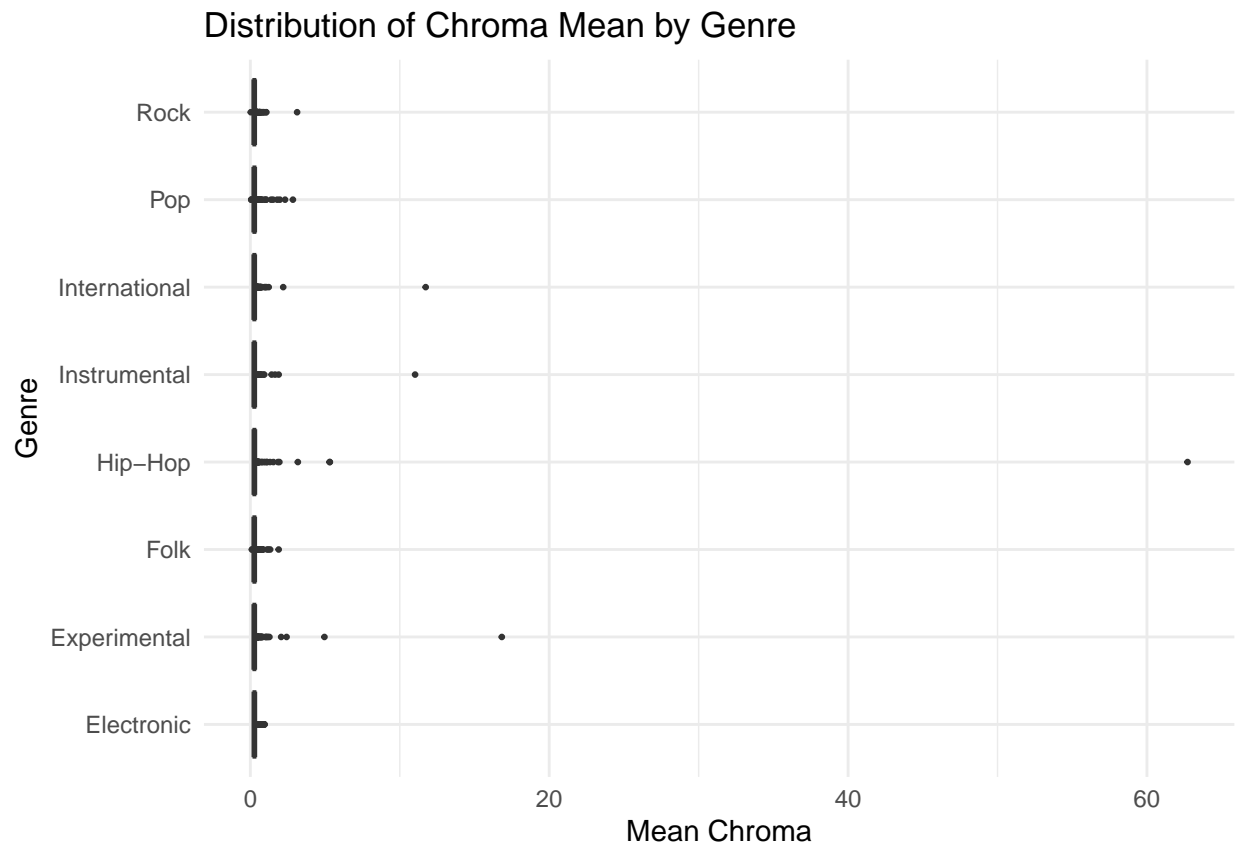
ggplot(combined_data, aes(x = genre_top, y = mfcc_mean)) +
  geom_boxplot(outlier.size = 0.5, fill = "lightblue") +
  coord_flip() +
  labs(title = "Distribution of MFCC Mean by Genre",
       x = "Genre", y = "Mean MFCC") +

```

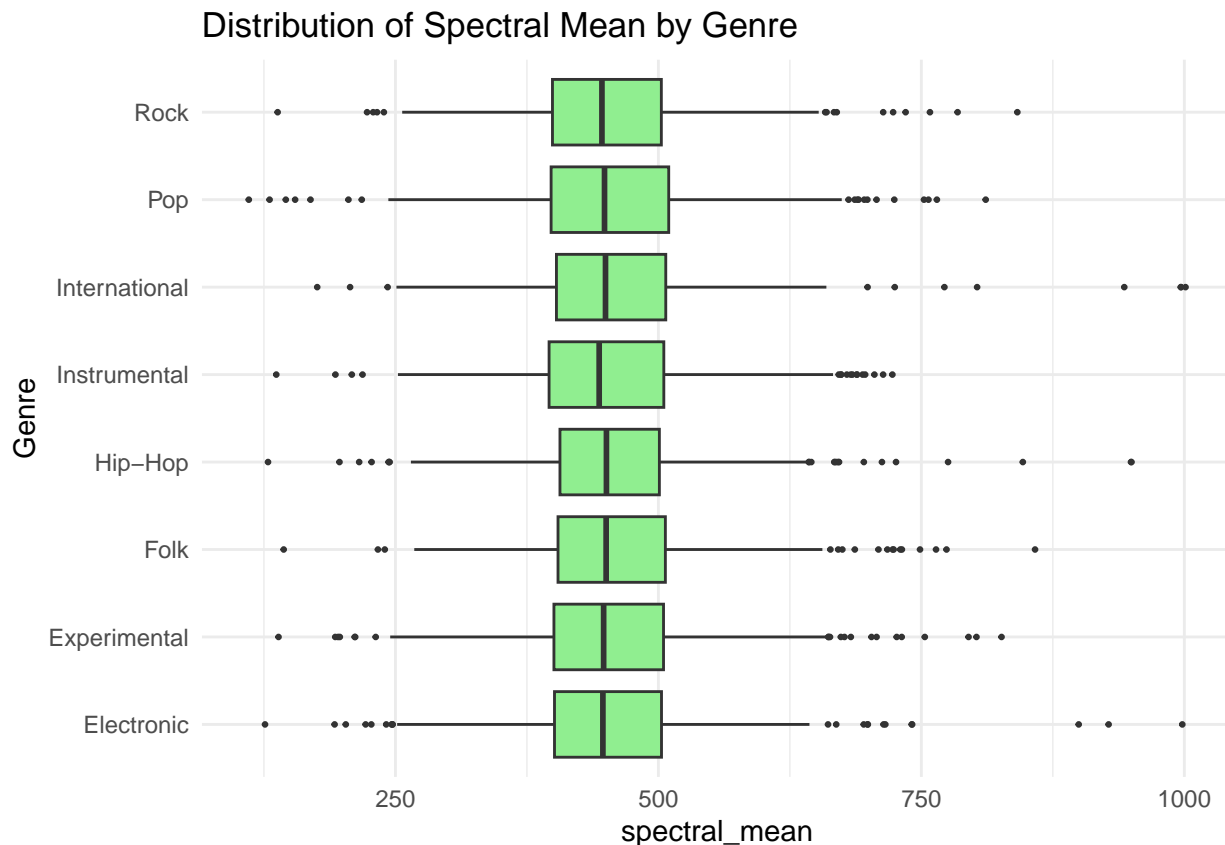
```
theme_minimal()
```



```
ggplot(combined_data, aes(x = genre_top, y = chroma_mean)) +  
  geom_boxplot(outlier.size = 0.5, fill = "lightgreen") +  
  coord_flip() +  
  labs(title = "Distribution of Chroma Mean by Genre",  
        x = "Genre", y = "Mean Chroma") +  
  theme_minimal()
```



```
ggplot(combined_data, aes(x = genre_top, y = spectral_mean)) +
  geom_boxplot(outlier.size = 0.5, fill = "lightgreen") +
  coord_flip() +
  labs(title = "Distribution of Spectral Mean by Genre",
        x = "Genre", y = "spectral_mean") +
  theme_minimal()
```



```
# Drop first row of features
features <- features[-1, ]
features$track_id <- as.integer(features$track_id)

# Merge features with cleaned tracks
combined_data <- inner_join(tracks_clean, features, by = "track_id")

# Convert relevant columns to numeric
mfcc_cols <- grep("^mfcc", names(combined_data), value = TRUE)
chroma_cols <- grep("^chroma", names(combined_data), value = TRUE)
spectral_cols <- grep("^spectral", names(combined_data), value = TRUE)

combined_data <- combined_data %>%
  mutate(across(all_of(c(mfcc_cols, chroma_cols, spectral_cols)), ~as.numeric(.)))

# Create model_data with scaled numeric features
numeric_features <- combined_data %>%
  select(-track_id, -subset, -genre_top) %>%
  select(where(is.numeric))


feature_matrix <- scale(numeric_features)
```

```
model_data <- data.frame(genre_top = combined_data$genre_top, feature_matrix)
```

```
set.seed(123)
train_index <- createDataPartition(model_data$genre_top, p = 0.8, list = FALSE)
train_set <- model_data[train_index, ]
test_set <- model_data[-train_index, ]
```

```
library(corrplot)
```

```
# Compute correlation matrix on a subset to avoid overload
corr_matrix <- cor(model_data[, 2:30]) # First 30 features
corrplot(corr_matrix, method = "color", type = "upper", tl.cex = 0.6)
```



```
#PCA for Visualization
```

```
library(FactoMineR)
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WB
```



```

# Run PCA on scaled features
pca_res <- prcomp(model_data[, -1], center = TRUE, scale. = TRUE)

# Plot first 2 PCs, colored by genre
fviz_pca_ind(pca_res,
  geom.ind = "point",
  habillage = model_data$genre_top,
  addEllipses = TRUE,
  palette = "Dark2",
  title = "PCA: Genre Clusters")

```



```

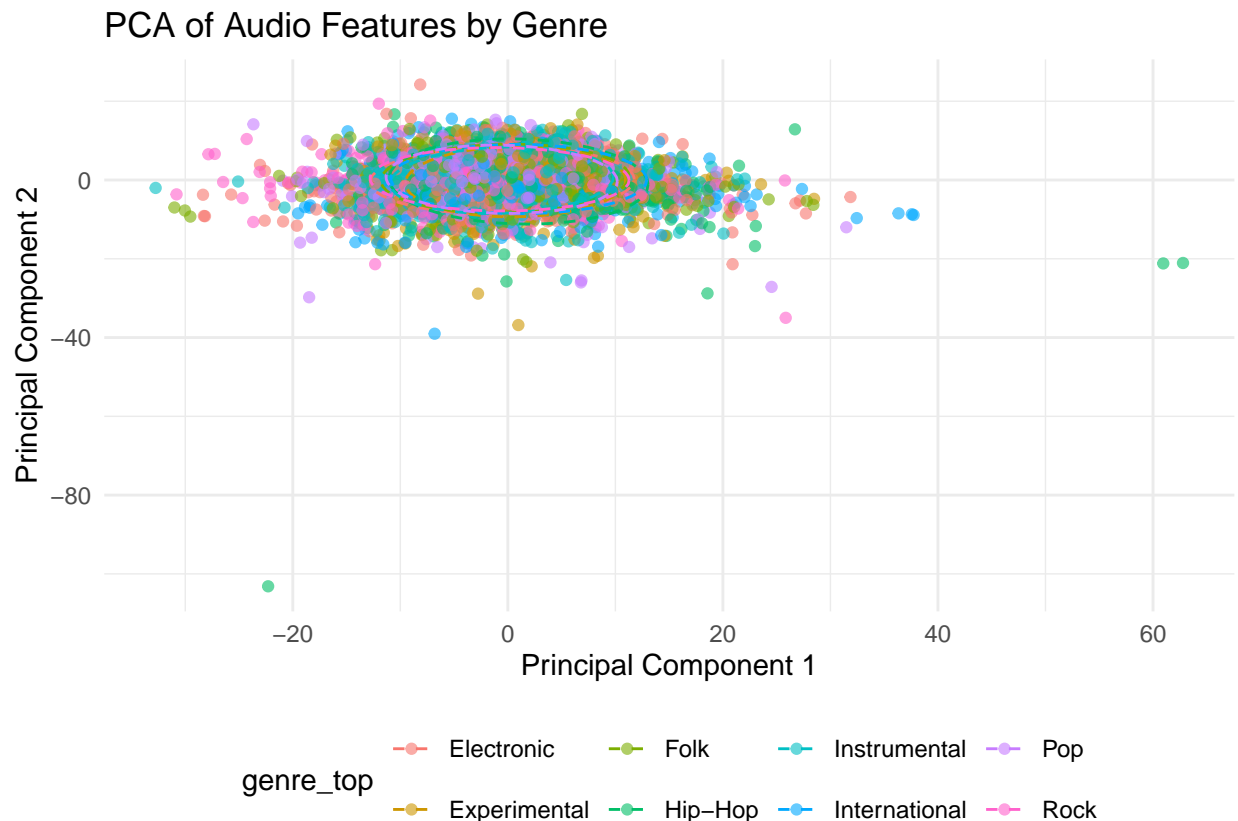
# Create pca_df
pca_result <- prcomp(model_data[, -1], center = TRUE, scale. = TRUE)
# Extract the first two principal components
pca_scores <- as.data.frame(pca_result$x[, 1:2])

# Add genre labels back in from model_data
pca_df <- cbind(pca_scores, genre_top = model_data$genre_top)

# PCA of Audio Features by Genre


```

```
ggplot(pca_df, aes(x = PC1, y = PC2, color = genre_top)) +
  geom_point(alpha = 0.6) +
  stat_ellipse(type = "norm", level = 0.68, linetype = "dashed") +
  labs(title = "PCA of Audio Features by Genre",
       x = "Principal Component 1",
       y = "Principal Component 2") +
  theme_minimal() +
  theme(legend.position = "bottom")
```



```
pca_variance <- pca_result$sdev^2
pca_prop_var <- pca_variance / sum(pca_variance)
qplot(y = pca_prop_var[1:10], x = 1:10, geom = "line") +
  labs(title = "Scree Plot", x = "Principal Component", y = "Proportion of Variance Explained")
```

```
## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```



1 Methods: Alexandria Simms

2 K Nearest Neighbors Model

```
knn_model <- train(
  genre_top ~ ., data = train_set,
  method = "knn",
  tuneLength = 5,
  trControl = trainControl(method = "cv", number = 5)
)
knn_preds <- predict(knn_model, newdata = test_set)
confusionMatrix(knn_preds, test_set$genre_top)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

## Prediction	Electronic	Experimental	Folk	Hip-Hop	Instrumental	International
## Electronic	24	19	13	11	19	12
## Experimental	22	26	15	24	24	23
## Folk	14	13	20	9	11	8
## Hip-Hop	12	11	13	28	12	16
## Instrumental	15	11	9	14	22	15
## International	18	19	16	19	14	33
## Pop	17	24	16	13	24	18
## Rock	11	13	9	10	18	11

```
##           Reference
```

```
## Prediction      Pop Rock
```

## Electronic	12	16
## Experimental	18	22
## Folk	10	12
## Hip-Hop	22	15
## Instrumental	16	14
## International	16	18
## Pop	37	12
## Rock	11	23

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.2006
```

```
##           95% CI : (0.1769, 0.2259)
```

```
##           No Information Rate : 0.1356
```

```
##           P-Value [Acc > NIR] : 3.153e-09
```

```
##
##          Kappa : 0.0854
##
##  McNemar's Test P-Value : 0.3524
##
## Statistics by Class:
##
##          Class: Electronic Class: Experimental Class: Folk
## Sensitivity          0.1805          0.19118      0.18018
## Specificity          0.8902          0.84017      0.91903
## Pos Pred Value       0.1905          0.14943      0.20619
## Neg Pred Value       0.8835          0.87613      0.90570
## Prevalence           0.1252          0.12806      0.10452
## Detection Rate       0.0226          0.02448      0.01883
## Detection Prevalence 0.1186          0.16384      0.09134
## Balanced Accuracy     0.5353          0.51567      0.54961
##
##          Class: Hip-Hop Class: Instrumental Class: International
## Sensitivity          0.21875          0.15278          0.24265
## Specificity          0.89186          0.89760          0.87041
## Pos Pred Value       0.21705          0.18966          0.21569
## Neg Pred Value       0.89282          0.87104          0.88669
## Prevalence           0.12053          0.13559          0.12806
## Detection Rate       0.02637          0.02072          0.03107
## Detection Prevalence 0.12147          0.10923          0.14407
## Balanced Accuracy     0.55531          0.52519          0.55653
##
##          Class: Pop Class: Rock
## Sensitivity          0.26056          0.17424
## Specificity          0.86522          0.91075
## Pos Pred Value       0.22981          0.21698
## Neg Pred Value       0.88346          0.88598
## Prevalence           0.13371          0.12429
## Detection Rate       0.03484          0.02166
## Detection Prevalence 0.15160          0.09981
## Balanced Accuracy     0.56289          0.54250
```

3 Random Forest Model

```
rf_model <- randomForest(
  genre_top ~ ., data = train_set,
  ntree = 100, importance = TRUE
)
rf_preds <- predict(rf_model, newdata = test_set)
```

```
confusionMatrix(rf_preds, test_set$genre_top)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

## Prediction	Electronic	Experimental	Folk	Hip-Hop	Instrumental	International
## Electronic	20	18	17	15	21	18
## Experimental	18	28	16	20	8	15
## Folk	6	6	2	2	8	6
## Hip-Hop	18	15	13	24	13	14
## Instrumental	20	21	21	18	40	23
## International	14	22	12	23	17	31
## Pop	29	14	22	19	24	16
## Rock	8	12	8	7	13	13

```
##           Reference
```

```
## Prediction      Pop Rock
```

## Electronic	17	26
## Experimental	16	17
## Folk	8	6
## Hip-Hop	11	6
## Instrumental	23	15
## International	16	18
## Pop	36	23
## Rock	15	21

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.1902
```

```
##           95% CI : (0.167, 0.2151)
```

```
##           No Information Rate : 0.1356
```

```
##           P-Value [Acc > NIR] : 4.396e-07
```

```
##
```

```
##           Kappa : 0.0716
```

```
##
```

```
##           McNemar's Test P-Value : 0.0001385
```

```
##
```

```
## Statistics by Class:
```

```
##
```

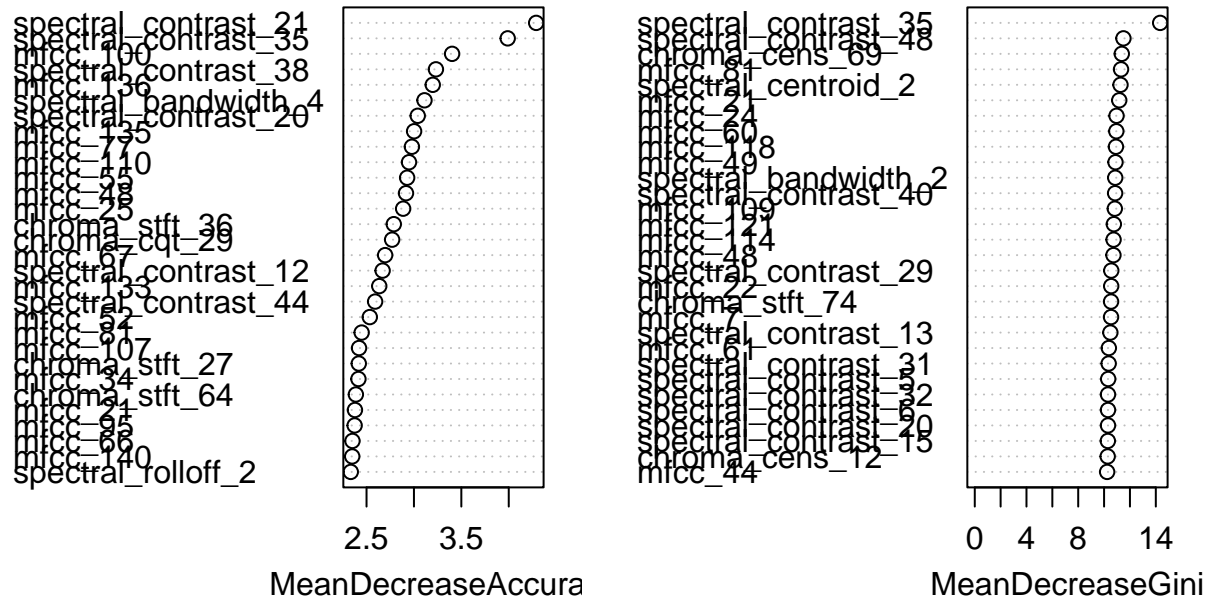
```
##           Class: Electronic Class: Experimental Class: Folk
```

## Sensitivity	0.15038	0.20588	0.018018
## Specificity	0.85791	0.88121	0.955836
## Pos Pred Value	0.13158	0.20290	0.045455
## Neg Pred Value	0.87582	0.88312	0.892927
## Prevalence	0.12524	0.12806	0.104520

## Detection Rate	0.01883	0.02637	0.001883
## Detection Prevalence	0.14313	0.12994	0.041431
## Balanced Accuracy	0.50414	0.54355	0.486927
##	Class: Hip-Hop	Class: Instrumental	Class: International
## Sensitivity	0.1875	0.27778	0.22794
## Specificity	0.9036	0.84641	0.86825
## Pos Pred Value	0.2105	0.22099	0.20261
## Neg Pred Value	0.8903	0.88195	0.88449
## Prevalence	0.1205	0.13559	0.12806
## Detection Rate	0.0226	0.03766	0.02919
## Detection Prevalence	0.1073	0.17043	0.14407
## Balanced Accuracy	0.5456	0.56209	0.54810
##	Class: Pop	Class: Rock	
## Sensitivity	0.2535	0.15909	
## Specificity	0.8402	0.91828	
## Pos Pred Value	0.1967	0.21649	
## Neg Pred Value	0.8794	0.88497	
## Prevalence	0.1337	0.12429	
## Detection Rate	0.0339	0.01977	
## Detection Prevalence	0.1723	0.09134	
## Balanced Accuracy	0.5469	0.53869	

```
varImpPlot(rf_model)
```

rf_model




4 Model Tuning

```
set.seed(123)
holdout_index <- createDataPartition(model_data$genre_top, p = 0.8, list = FALSE)
training_set <- model_data[holdout_index, ]
holdout_set <- model_data[-holdout_index, ]

training_set$genre_top <- as.factor(training_set$genre_top)
holdout_set$genre_top <- as.factor(holdout_set$genre_top)

# KNN Tuning
tuned_knn <- train(
  genre_top ~ ., data = training_set,
  method = "knn", tuneLength = 10,
  trControl = trainControl(method = "cv", number = 5)
)
plot(tuned_knn)
```

```
knn_final_preds <- predict(tuned_knn, newdata = holdout_set)
confusionMatrix(knn_final_preds, holdout_set$genre_top)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

## Prediction	Electronic	Experimental	Folk	Hip-Hop	Instrumental	International
## Electronic	25	13	12	11	21	15
## Experimental	17	30	17	21	25	23
## Folk	14	14	18	7	10	6
## Hip-Hop	14	12	14	29	11	20
## Instrumental	14	9	11	13	22	11
## International	15	20	18	20	17	32
## Pop	20	22	14	17	23	16
## Rock	14	16	7	10	15	13

```
##           Reference
```

```
## Prediction      Pop Rock
```

## Electronic	14	12
## Experimental	18	23
## Folk	13	13
## Hip-Hop	22	13

```


## Instrumental    17    14
## International  18    24
## Pop            30    12
## Rock           10    21
##
## Overall Statistics
##
## Accuracy : 0.1949
## 95% CI : (0.1715, 0.22)
## No Information Rate : 0.1356
## P-Value [Acc > NIR] : 5.083e-08
##
## Kappa : 0.079
##
## McNemar's Test P-Value : 0.2355
##
## Statistics by Class:
##
## Class: Electronic Class: Experimental Class: Folk
## Sensitivity          0.18797          0.22059          0.16216
## Specificity          0.89451          0.84449          0.91903
## Pos Pred Value       0.20325          0.17241          0.18947
## Neg Pred Value       0.88498          0.88063          0.90383
## Prevalence           0.12524          0.12806          0.10452
## Detection Rate       0.02354          0.02825          0.01695
## Detection Prevalence 0.11582          0.16384          0.08945
## Balanced Accuracy     0.54124          0.53254          0.54060
##
## Class: Hip-Hop Class: Instrumental Class: International
## Sensitivity          0.22656          0.15278          0.23529
## Specificity          0.88651          0.90305          0.85745
## Pos Pred Value       0.21481          0.19820          0.19512
## Neg Pred Value       0.89320          0.87171          0.88419
## Prevalence           0.12053          0.13559          0.12806
## Detection Rate       0.02731          0.02072          0.03013
## Detection Prevalence 0.12712          0.10452          0.15443
## Balanced Accuracy     0.55654          0.52791          0.54637
##
## Class: Pop Class: Rock
## Sensitivity          0.21127          0.15909
## Specificity          0.86522          0.90860
## Pos Pred Value       0.19481          0.19811
## Neg Pred Value       0.87665          0.88389
## Prevalence           0.13371          0.12429
## Detection Rate       0.02825          0.01977
## Detection Prevalence 0.14501          0.09981
## Balanced Accuracy     0.53824          0.53385

```

```

# RF Tuning
rf_grid <- expand.grid(mtry = c(10, 20, 30, 40, 50))
rf_tuned <- train(
  genre_top ~ ., data = training_set,
  method = "rf",
  trControl = trainControl(method = "cv", number = 5),
  tuneGrid = rf_grid,
  importance = TRUE
)
plot(rf_tuned)

```



```

rf_final_preds <- predict(rf_tuned, newdata = holdout_set)
cm_rf <- confusionMatrix(rf_final_preds, holdout_set$genre_top)
print(cm_rf)

```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      Electronic Experimental Folk Hip-Hop Instrumental International
```

```
## Electronic      29           9    10     11           11           12
```

```

##      Experimental      11      35  16      20      16      21
##      Folk              4       1   9       1       0       2
##      Hip-Hop           13      13   9      26       7      10
##      Instrumental      20      24  23      18      43      26
##      International     16      28  13      22      24      34
##      Pop               29      18  20      23      27      22
##      Rock              11       8  11       7      16       9
##
##      Reference
## Prediction      Pop Rock
##      Electronic   17   11
##      Experimental 15   29
##      Folk          3    1
##      Hip-Hop      10    8
##      Instrumental 25   22
##      International 22   14
##      Pop          41   22
##      Rock          9   25
##
## Overall Statistics
##
##      Accuracy : 0.2279
##      95% CI : (0.203, 0.2543)
##      No Information Rate : 0.1356
##      P-Value [Acc > NIR] : 2.713e-16
##
##      Kappa : 0.1137
##
##      McNemar's Test P-Value : 3.070e-13
##
## Statistics by Class:
##
##      Class: Electronic Class: Experimental Class: Folk
## Sensitivity      0.21805      0.25735      0.081081
## Specificity      0.91281      0.86177      0.987382
## Pos Pred Value   0.26364      0.21472      0.428571
## Neg Pred Value   0.89076      0.88765      0.902017
## Prevalence       0.12524      0.12806      0.104520
## Detection Rate   0.02731      0.03296      0.008475
## Detection Prevalence 0.10358      0.15348      0.019774
## Balanced Accuracy 0.56543      0.55956      0.534231
##
##      Class: Hip-Hop Class: Instrumental Class: International
## Sensitivity      0.20312      0.29861      0.25000
## Specificity      0.92505      0.82789      0.84989
## Pos Pred Value   0.27083      0.21393      0.19653
## Neg Pred Value   0.89441      0.88269      0.88526

```

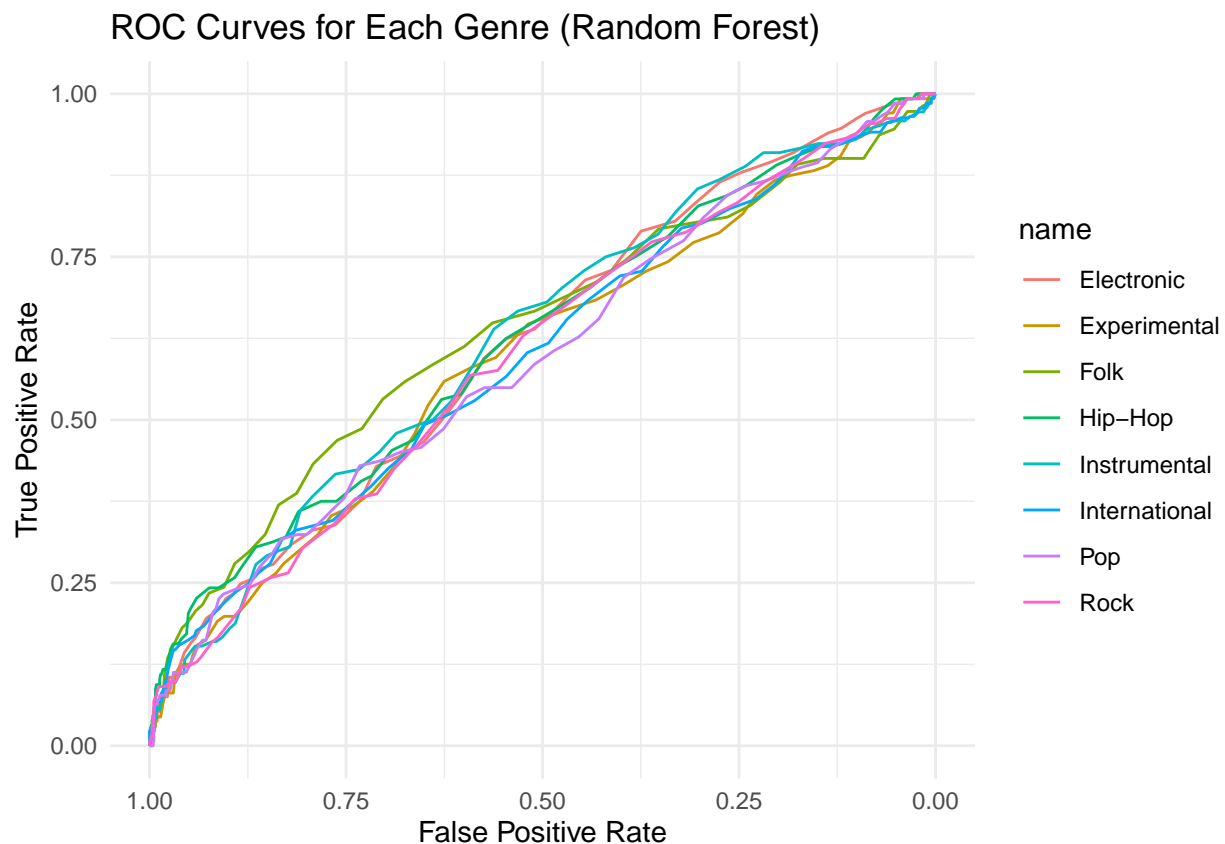
## Prevalence	0.12053	0.13559	0.12806
## Detection Rate	0.02448	0.04049	0.03202
## Detection Prevalence	0.09040	0.18927	0.16290
## Balanced Accuracy	0.56409	0.56325	0.54995
##	Class: Pop	Class: Rock	
## Sensitivity	0.28873	0.18939	
## Specificity	0.82500	0.92366	
## Pos Pred Value	0.20297	0.26042	
## Neg Pred Value	0.88256	0.88923	
## Prevalence	0.13371	0.12429	
## Detection Rate	0.03861	0.02354	
## Detection Prevalence	0.19021	0.09040	
## Balanced Accuracy	0.55687	0.55652	

5 Results

```
rf_probabilities <- predict(rf_tuned, newdata = holdout_set, type = "prob")
roc_list <- list()
true_labels <- holdout_set$genre_top
classes <- colnames(rf_probabilities)

for (class in classes) {
  binary_labels <- ifelse(true_labels == class, 1, 0)
  prob <- rf_probabilities[[class]]
  roc_obj <- roc(binary_labels, prob, levels = c(0, 1), direction = "<")
  roc_list[[class]] <- roc_obj
}

ggroc(roc_list) +
  labs(title = "ROC Curves for Each Genre (Random Forest)",
       x = "False Positive Rate", y = "True Positive Rate") +
  theme_minimal()
```



```
sapply(roc_list, auc)
```

```
##      Electronic  Experimental      Folk      Hip-Hop  Instrumental
##      0.6161569      0.5954096      0.6364140      0.6216667      0.6245991
## International      Pop      Rock
##      0.5979744      0.5958550      0.5995194
```

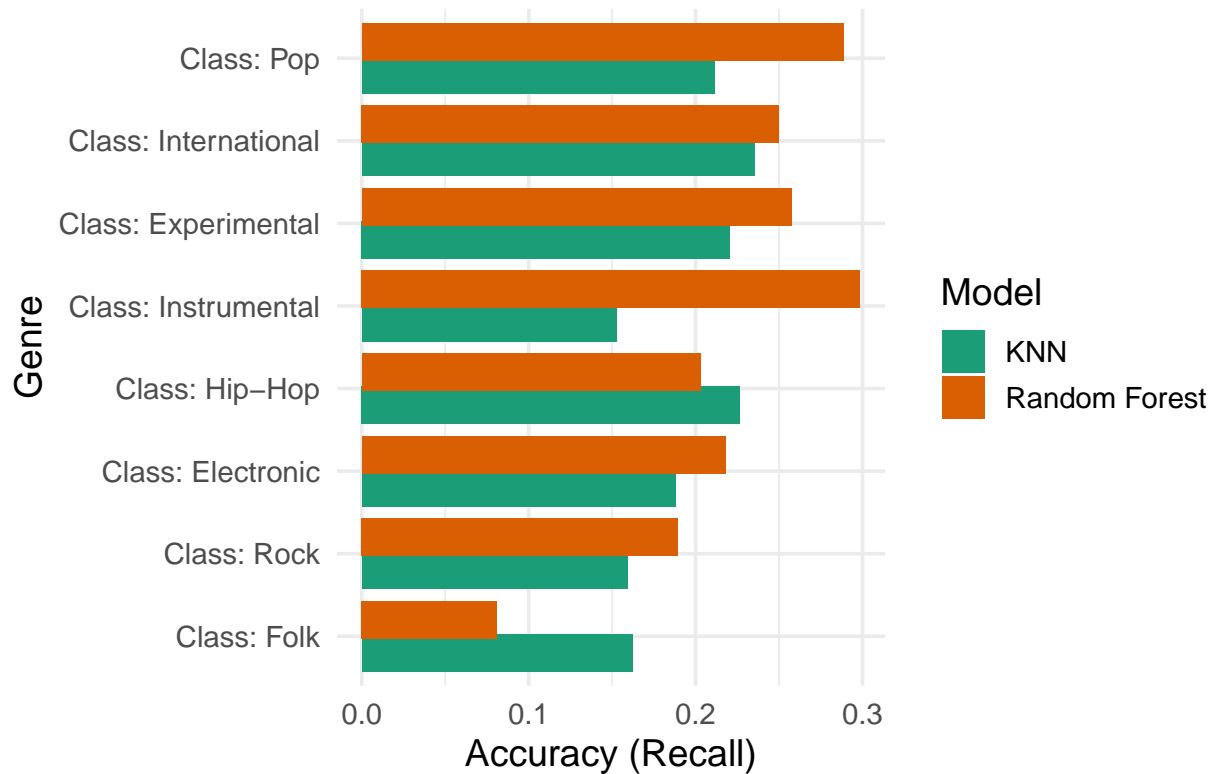
```
cm_knn <- confusionMatrix(knn_final_preds, holdout_set$genre_top)
knn_accuracy <- as.data.frame(cm_knn$byClass) %>%
  mutate(Genre = rownames(.), Model = "KNN") %>%
  select(Genre, Sensitivity, Model)
```

```
rf_accuracy <- as.data.frame(cm_rf$byClass) %>%
  mutate(Genre = rownames(.), Model = "Random Forest") %>%
  select(Genre, Sensitivity, Model)
```

```
combined_accuracy <- bind_rows(knn_accuracy, rf_accuracy)
```

```
ggplot(combined_accuracy, aes(x = reorder(Genre, Sensitivity), y = Sensitivity, fill = Model)) +
  geom_bar(stat = "identity", position = position_dodge(width = 0.8)) +
  coord_flip() +
  labs(title = "Genre-wise Accuracy Comparison: KNN vs Random Forest",
       x = "Genre", y = "Accuracy (Recall)") +
  theme_minimal(base_size = 14) +
  scale_fill_brewer(palette = "Dark2")
```

Genre-wise Accuracy Comparison: KNN vs F



```
# Get variable importance from rf_tuned
rf_importance <- varImp(rf_tuned, scale = TRUE)

# Calculate average importance across genres for each feature
rf_importance_overall <- rf_importance$importance %>%
  rowMeans() %>%
  sort(decreasing = TRUE)

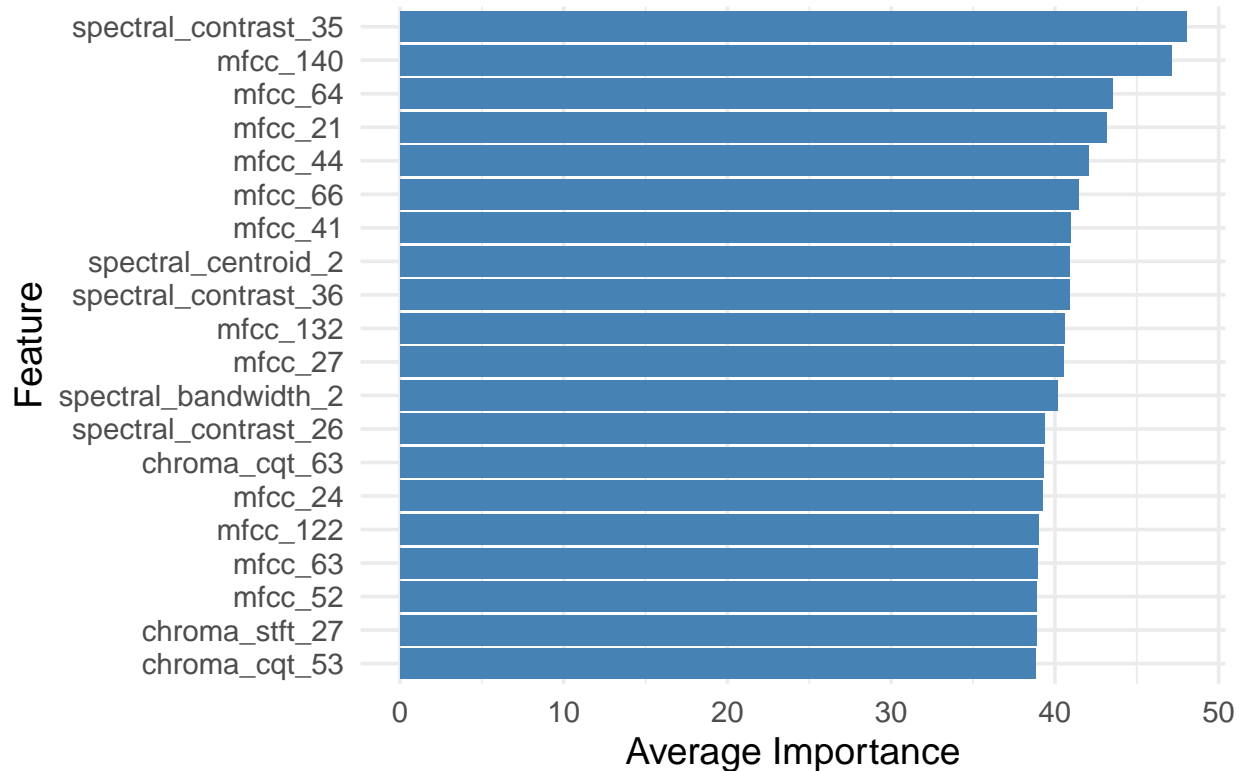
# Convert to data frame
rf_top_features <- data.frame(
  Feature = names(rf_importance_overall),
  Importance = rf_importance_overall
)

# Plot Top 20 features
rf_top_features %>%
  slice_max(order_by = Importance, n = 20) %>%
  ggplot(aes(x = reorder(Feature, Importance), y = Importance)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  coord_flip() +
  labs(title = "Top 20 Most Informative Features (Random Forest)",
```



```
x = "Feature", y = "Average Importance") +  
theme_minimal(base_size = 14)
```

Top 20 Most Informative Features (Random



```
# Genre Frequency in Train vs Test vs Holdout sets
```

```
library(ggplot2)  
library(dplyr)
```

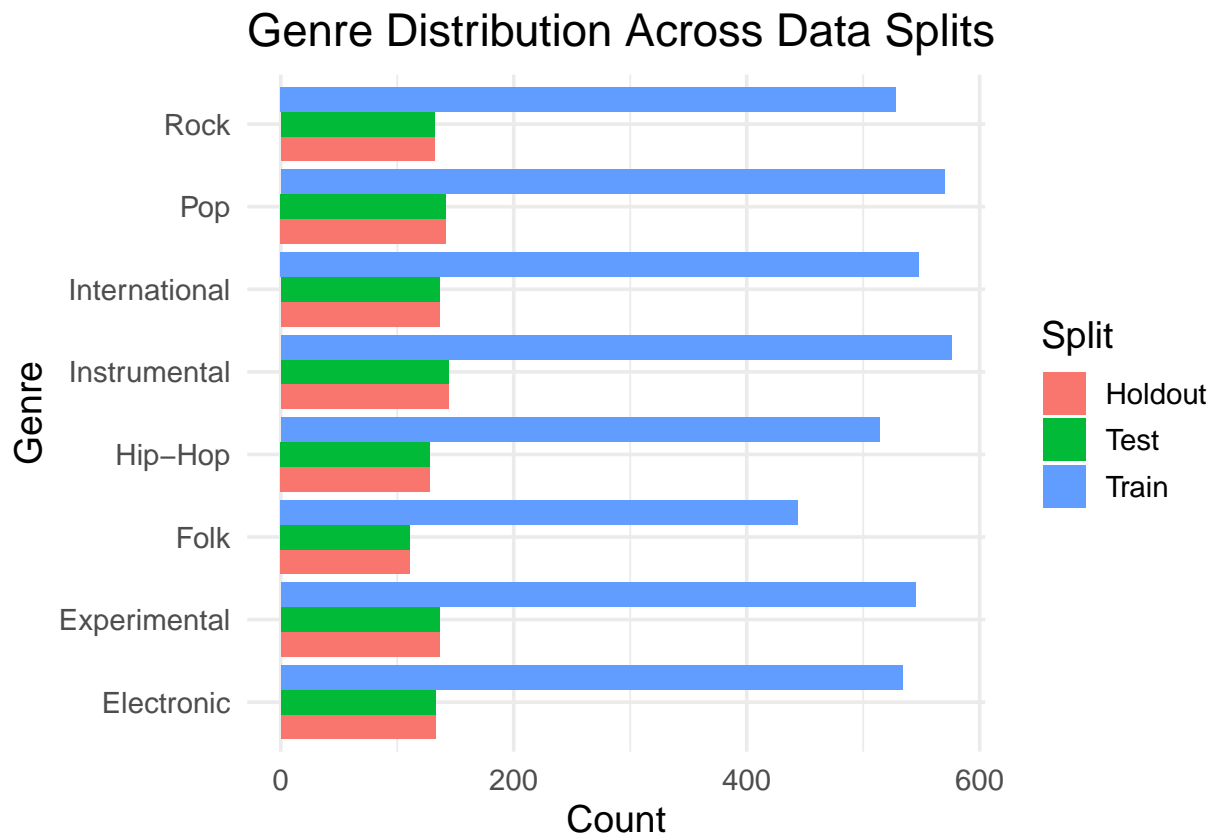
```
# Add labels for dataset splits
```

```
genre_distribution <- bind_rows(  
  train_set %>% mutate(Split = "Train"),  
  test_set %>% mutate(Split = "Test"),  
  holdout_set %>% mutate(Split = "Holdout")  
)
```

```
# Plot distribution
```

```
ggplot(genre_distribution, aes(x = genre_top, fill = Split)) +  
  geom_bar(position = "dodge") +  
  labs(title = "Genre Distribution Across Data Splits",  
       x = "Genre", y = "Count") +
```

```
coord_flip() +  
theme_minimal(base_size = 14)
```



```
save(train_set, test_set, rf_model, knn_model, file = "models_and_data.RData")
```